

Recommender Systems: Part 2

Web Science Lecture

22 March 2020

Sagnik Ray Choudhury

src@di.ku.dk

<http://sagnik.in>

UNIVERSITY OF COPENHAGEN



Organization

- Part 1: Collaborative Filtering: Better Representation Through Dimensionality Reduction: PCA and SVD
- Part2: Latent factor models: more discussions.
- Part3: Intro to DNNs for RecSys
- Acknowledgements:
 - Charu C Agarwal ([Recommender Systems, The Textbook](#))
 - Jon Shlens ([PCA](#)).
 - Jeremy Kun ([SVD](#))

Announcement

- On **March 15** you will receive a link for the course evaluation in your KU e-mailbox.
- The deadline to complete the course evaluation is **March 28**.
- Please remember to fill out the course evaluation, even if you do not have anything wrong to report!
- Your feedback is very important to us, so we can understand how to improve this course for the next years.

Recommender Systems: Recap

- Problem definition: Given a utility matrix ($U \times N$, $U = |\text{Users}|$, $N = |\text{Items}|$) with some values unspecified, 1. predict the unspecified values or 2. Rank the items according to the relevance of the users.
 - 2 comes naturally from 1.
- Evaluation methods for recommender systems
 - RMSE (evaluate 1)
 - Precision, Recall, $F1@k$. (evaluate 2)
 - Average precision @K, $MAP@K$.
 - Cumulative gain, DCG and NDCG.
- Recommender system algorithms
 - Content based
 - Collaborative filtering: User-User and Item-Item.
 - Latent Factor Models

Recommender Systems: Recap

- Content-based methods:
 - Start by representing an item in a feature space.
 - Represent user by combining the item she rates -> user is represented in the item feature space.
 - For a new item, the score is calculated by a distance metric between the user and the items.
 - Feature space is manually engineered: **hard to create**.
- Collaborative filtering:
 1. For a target entry (u, i) determine the most similar rows/columns of the ratings matrix with the use of the cosine coefficient between rows/columns. For user-based methods rows are used, whereas for item-based methods, columns are used.
 2. Predict the target entry (u, i) using a weighted combination of the ratings in the most similar rows/columns determined in the first step.
 - Item (user) feature space is the users (items): this is sparse.

Latent Factor Models: Sparsity to Compactness

- Consider a utility matrix $U \times I$ (number of users \times number of items).
- Collaborative Filtering problems:
 - Feature space for item/user is sparse: lots of redundancy.
 - But complexity is $\mathcal{O}(N^2U)$ (item-item) . Reducing N to d ($d \ll N$) would be a huge improvement.
 - Side note: This happens in word vectors as well.
- Possible solution (for user-user, works the same way for items as well):
 - Reduce the dimensionality to d where $d \ll U$ and I .
 - This will produce a new representation of the data.
 - Use this to represent each user and then repeat the usual steps.
- What dimensionality reduction technique to use?
 - PCA.
 - Singular value decomposition (SVD).

Dimensionality Reduction: Refresher

- Users are represented in item space and items are represented in user space.
- It's a vector space: it has a set of *basis vectors*. The basis vectors *span* that space.
- Basis vectors are linearly independent. A span of basis vectors is a set of vectors that can be expressed as a linear combination of these basis vectors.

1. Ratings between -1 and 1.
2. No missing values

NERO	Julius Caesar	Patriot	Sleepless Seattle	Pretty Woman	Casablanca
1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	1	1	1
-1	-1	-1	1	1	1
-1	-1	-1	1	1	1
-1	-1	-1	1	1	1

- There's a correlation in the data.
- There is one vector that represents user 1-3, 1 vector that can represent user 5-7, and one vector that can represent user 4. So obviously, we can represent each user by at least them.

1	1	1	0	0	0
1	1	1	1	1	1
-1	-1	-1	1	1	1

1	0	0
1	0	0
1	0	0
0	1	0
0	0	1
0	0	1
0	0	1

New (not quite) basis vectors

- What is the benefit of that?
- Dimensionality is reduced to 3.
- Can we do a little better?

Is This the Best We Can Do?

- (Row/column) Rank of a matrix is the number of linearly independent vectors (they are the same, why?).
- A data matrix \mathbf{X} [$m \times n$] can be represented in a k dimensional subspace where $\text{rank}(\mathbf{X}) = k < (m, n)$
- $\mathbf{X} = \mathbf{U} \times \mathbf{V}^T$ (\mathbf{U} = new data matrix, \mathbf{V} new basis vectors, with $\mathbf{U} = [m \times k]$ and $\mathbf{V} = [n \times k]$)
- If $\text{rank}(\mathbf{X}) = k$, that decomposition always exists.

This matrix has rank 2. Verify.

1	0
1	0
1	0
1	1
-1	1
-1	1
-1	1

NERO	Julius Caesar	Patriot	Sleepless Seattle	Pretty Woman	Casablanca
1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	1	1	1
-1	-1	-1	1	1	1
-1	-1	-1	1	1	1
-1	-1	-1	1	1	1

1. Ratings between -1 and 1.
2. No missing values

Here, the basis vectors are changed to $[1, 1, 1, 0, 0, 0]$ and $[0, 0, 0, 1, 1, 1]$. That is a reduction from 6 to 2. Now, we need exactly 2 dimensions to represent the data.

1	1	1	0	0	0
0	0	0	1	1	1

Is This the Best We Can Do?: Contd.

1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	1	1	1
-1	-1	1	1	1	1
-1	-1	-1	1	1	1
-1	-1	-1	1	1	1

Previously, we had rank 2. This means, we could represent the data with 2 basis vectors. Now, we have a matrix of rank 3, so we will need at least 3 such vectors. Equivalently, We can not completely represent this with 2 dimensions, so there will be some error.

1	0
1	0
1	0
1	1
-1	1
-1	1
-1	1

U

1	1	1	0	0	0
0	0	1	1	1	1

V^T

1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	0	0	0
1	1	2	1	1	1
-1	-1	0	1	1	1
-1	-1	0	1	1	1
-1	-1	0	1	1	1

U x V^T

- Let's assume we are okay with losing some information.
- But **how do we find out these matrices U and V?**
- If the data matrix had rank k , not hard to reduce the ranks. But what happens when we want to *approximate* k ?

PCA: we will find out a projection matrix **P** such that **XP = Y**.

Original formulation was **X = UV^T**.

Not hard to see that **(V^T)⁻¹** would be **P**.

PCA

- We are projecting the data to a new space with a *linear* transformation.
- This is the same as changing the basis vectors (these new basis vectors spans the new space, and we are okay with some error).
- What are the properties we would want in such a transformation?
 - We want to minimize the reconstruction error.
 - We want to project the data in the directions of highest variance.
 - User 1-3 (5-7) are similar to each other. But User 3,4,5 are dissimilar.
 - If there's one dimension, you would want to have 3 numbers each significantly different.
 - As you can see, that is not possible in this data. You need at least 2 dimensions to express that.
- Turns out, we can choose a projection matrix that satisfies both this criteria.
 - We showed how in the last class, but we will do that here again.

1	0
1	0
1	0
1	1
-1	1
-1	1
-1	1

Covariance

- Variance: how do the values of a variable vary across itself [X takes values $(x_1 \dots x_n)$, bars indicates mean]
- Two variables, X $(x_1 \dots x_n)$ and Y $(y_1 \dots y_n)$, how do the values vary across themselves?

$$s^2 = \frac{\sum (x_i - \bar{x})^2}{n - 1}$$

$$\text{cov}_{x,y} = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{N - 1}$$

$$\text{Cov}(X_i, X_j) = \frac{\sum_{k=1}^m (X_{ik} - \bar{X}_i)(X_{jk} - \bar{X}_j)}{m - 1}$$

Data matrix

$$\mathbf{M} = \begin{bmatrix} X_{11} & X_{12} & \dots & X_{1n} \\ X_{21} & X_{22} & \dots & X_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ X_{m1} & X_{m2} & \dots & X_{mn} \end{bmatrix}$$

Covariance Matrix

$$\begin{bmatrix} \text{Var}(X_1) & \text{Cov}(X_1, X_2) & \dots & \text{Cov}(X_1, X_n) \\ \text{Cov}(X_2, X_1) & \text{Var}(X_2) & \dots & \text{Cov}(X_2, X_n) \\ \vdots & \vdots & \ddots & \vdots \\ \text{Cov}(X_n, X_1) & \text{Cov}(X_n, X_2) & \dots & \text{Var}(X_n) \end{bmatrix}$$

- Each feature is a random variable that is represented in a column (called X_i).
- Each X_i takes values X_{i1} to X_{im} .
- There are n such columns (random variables), X_1 to X_n .

- $\text{Cov}(X_i, X_j) = \text{Cov}(X_j, X_i)$ (verify) => The matrix is symmetric.
- If you subtract the column means from the column values, $\text{Cov}(\mathbf{X}) = (\mathbf{X}^T \mathbf{X}) / m - 1$ (Under the assumption that the rows are data points and columns are features.).
- We will just take the numerator from now on because the denominator does not depend on \mathbf{X} values.

What is the Relationship with PCA?

- If two vectors (X_i, X_j) are highly correlated (think one as the multiple of the other), the covariance is high.
 - Also, if we had one of them as the basis vector, would we need the other?
- If the correlation is low (think two perpendicular vectors) the covariance is low (0 when the vectors are perpendicular to each other).
- PCA definition: $\mathbf{XP} = \mathbf{Y}$
- In the covariance matrix of \mathbf{Y} (\mathbf{C}_Y), we would want the diagonal entries (variances) to be high, and the non diagonal entries (covariances) to be low.
- In other words, we would want \mathbf{C}_Y to be a *diagonal matrix*.
- What is the projection matrix \mathbf{P} that makes this happen?

Finding the Proper P

- Eigenvector of a matrix **A** is a vector **X** such that: $\mathbf{AX} = c\mathbf{X}$ (c is a constant).
 - What does that mean, intuitively?
 - A matrix is a rotation + scaling transformation on a vector. Eigenvectors are those vectors which only gets scaled by this operation.
 - We only consider the *normalized* ones, because if **E** is an eigenvector, so is $d\mathbf{E}$ ($d=\text{constant}$)
- A symmetric matrix is diagonalized by an orthonormal matrix of its eigenvectors.

Let's assume this is true for now, I will upload the proof separately

 - If there is a symmetric matrix **A** with associated eigenvectors $\{\mathbf{e}_1, \mathbf{e}_2, \dots, \mathbf{e}_n\}$, and **E** is a matrix such that the i^{th} column of **E** is \mathbf{e}_i , then there exists a diagonal matrix **D** such that $\mathbf{A} = \mathbf{EDE}^T$.
- We are looking for a **P** that makes **C_Y** (Covariance matrix of **Y**) diagonal.
- The trick: Let us select **P** to be the matrix where each column is the eigen vector of **C_X**.

Finding the Proper P, Contd.

Unit vectors, product of each other 0

- If there is a symmetric matrix **A** with associated *orthonormal* eigenvectors $\{e_1, e_2, \dots, e_n\}$ (they are always orthogonal if A is symmetric, *why?*), and **E** is a matrix such that the i^{th} column of E is e_i , then there exists a diagonal matrix **D** such that **A** = **EDE^T**.
- We are looking for a **P** that makes **C_Y** (Covariance matrix of **Y**) diagonal.
- The trick: Let us select **P** to be the matrix where each column is the orthonormal eigen vectors of **C_X**. Then we can write **C_X** as **PDP^T**.

Not considering denominators, hence approximate

$$\mathbf{C}_Y = \text{Cov}(\mathbf{Y}) = \text{Cov}(\mathbf{XP}) \sim (\mathbf{XP})^T(\mathbf{XP}) = \mathbf{P}^T\mathbf{X}^T\mathbf{XP} = \mathbf{P}^T\mathbf{C}_X\mathbf{P}$$

If I make P according to the last condition, we can write **C_X** = **PDP^T**. Replace **C_X** with this.

C_Y = **P^T (PDP^T) P**. Now, **P** is orthonormal, therefore, **P^TP** = **I**. Therefore, **C_Y** = **D** (which we wanted to do).

PCA, Finally.

- We would call the columns of \mathbf{P} (the projection vectors) the principal components.
 - This is obtained by diagonalizing \mathbf{C}_Y : this could be done in multiple ways.
 - PCA is one such diagonalization where each projection vector is orthonormal.
 - We also have a way to judge the importance of these dimensions (by the variance) .
 - Dimensions
 - Data (\mathbf{X}) dimensions: $m \times n$
 - Cov matrix \mathbf{C}_X : $n \times n$
 - $\text{Rank}(\mathbf{C}_X) = r \Rightarrow r$ orthonormal eigenvectors (which will come to be the principal components). You will choose the top k (The ones with the highest variance). This will make the projection matrix \mathbf{P} .
- We have found a way of dimensionality reduction that not only removes redundancy but also projects the data in an interesting subspace (variance maximization and covariance minimization).
 - The principal components are orthonormal: which helps to cast it as the problem of finding eigenvectors of the covariance matrix.
 - Still, computing that covariance matrix is time consuming: matrix multiplication (what is the complexity?) and over/underflow.

Singular Value Decomposition: Concepts

- $\mathbf{X} = m \times n$.
- v_1, v_2, \dots, v_r is the set of orthonormal eigenvectors of $\mathbf{X}^T \mathbf{X}$ with eigenvalues (e_1, e_2, \dots, e_r) .
 - $(\mathbf{X}^T \mathbf{X})v_i = e_i v_i$, v_i has a dimension of $n \times 1$;
 - As before, $\mathbf{X}^T \mathbf{X}$ is symmetric with rank r , therefore has r real eigenvalues and associated eigenvectors.
 - Let's define $s_i = \sqrt{e_i}$. We will call them singular values.
- Can you see that $\mathbf{X}v_i$ (dimension $m \times 1$) and $\mathbf{X}v_j$ are orthogonal (for any i, j)?
 - It turns out we can get a set $\{u_1, \dots, u_r\}$ of *orthonormal* vectors such that $\mathbf{X}v_i = s_i u_i$
- This is interesting.
 - $\{v_i\}$ and $\{u_i\}$ are orthonormal sets of vectors that span the column space **and** the row space.
 - For PCA, we got a set of orthonormal vectors that spanned *one* of row/ column space (depending on how you represent the data).

SVD, Construction.

- $\{v_i\}$ and $\{u_i\}$ are orthonormal sets of vectors that span the column space **and** the row space.
- Both sets have cardinality r where $\text{rank}(\mathbf{X}^T \mathbf{X}) = r$
- Let's build the matrices: $\mathbf{X}v_i = s_i u_i$ or, $\mathbf{X}v_i = u_i s_i$.

$s_1 = 1, s_2 = 1$

1	2	3
4	5	6
7	8	9
10	11	12

$\mathbf{X} (4 \times 3)$

1	2
3	4
5	6

$v_1 \quad v_2$

22	28
49	64
76	100
103	136

$u_1 \quad u_2$

1	2
3	4
5	6

$\mathbf{V} (3 \times 2)$

22	28
49	64
76	100
103	136

$\mathbf{U} (4 \times 2)$

These are not orthonormal vectors, here just to explain the construction!

- What happens when $s_1 = 5, s_2 = 3$?
- We add a diagonal matrix Σ
- In the diagonal matrix, the values will be in decreasing order (this would require u and v columns to be rearranged appropriately), and finally, $\mathbf{XV} = \mathbf{U} \Sigma$.
- Multiplying both side by \mathbf{V}^{-1} , $\mathbf{X} = \mathbf{U} \Sigma \mathbf{V}^{-1} \Rightarrow \mathbf{X} = \mathbf{U} \Sigma \mathbf{V}^T$. ($\mathbf{V}^T = \mathbf{V}^{-1}$ because V is orthonormal).

By column stacking, we can write these equations more concisely!, or $\mathbf{XV} = \mathbf{U}$

This is the SVD formulation.

SVD: Implications

- We know that a data matrix \mathbf{X} ($m \times n$) can be written as $\mathbf{X} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$, where:
 - \mathbf{U} is a set of orthonormal basis vectors ($\mathbf{U} \Rightarrow m \times r$).
 - \mathbf{V} is another set of orthonormal basis vectors ($\mathbf{V} \Rightarrow n \times r$).
 - $\mathbf{\Sigma}$ is a diagonal matrix, of some sort.
- What is really happening?

	Aisha	Bob	Chandrika
Up	2	5	3
Skyfall	1	2	1
Thor	4	1	1
Amelie	3	5	2
Snatch	5	3	1
Casablanca	4	5	5
Bridesmaids	2	4	2
Grease	2	2	5

SVD is a method to do this so that these two discoveries happen simultaneously.
Moreover, you can choose a rank k approximation of the data, by choosing $r=k$

- Some person rates a movie. It's a complex process. We have no idea about the process, but we have data that says 3 people rated 8 movies.
- With this data we will represent movies in a 3d space, and people in an 8d space.
- New movie comes in, we would want to represent it as a linear combination of existing movies and vice versa.
- We want to discover a special list of vectors $v_1 \dots v_8$ to do just that (we did something similar in PCA). The Same goes for the people $p_1, p_2 \dots p_3$.

SVD and PCA

- We don't have time to discuss how SVD is implemented.
 - In the QA session, we will go through a very basic solution.
- For now, let's assume very efficient algorithms exist for SVD.
- Can we use SVD to compute PCA?
- We know the principal components of \mathbf{X} are the eigen vectors of \mathbf{C}_X .
- Since \mathbf{C}_X is symmetric, we can also write $\mathbf{C}_X = \mathbf{E}\mathbf{D}\mathbf{E}^T$, where \mathbf{E} is the eigenvectors of \mathbf{C}_X .
- $\mathbf{C}_X = \mathbf{X}^T\mathbf{X}/n-1$. According to SVD, $\mathbf{X} = \mathbf{U}\mathbf{S}\mathbf{V}^T$.
- $\mathbf{C}_X = (\mathbf{U}\mathbf{S}\mathbf{V}^T)^T(\mathbf{U}\mathbf{S}\mathbf{V}^T)/n-1 = \mathbf{V}\mathbf{S}^T\mathbf{U}^T\mathbf{U}\mathbf{S}\mathbf{V}^T/n-1 = \mathbf{V}(\mathbf{S}^2/n-1)\mathbf{V}^T$? (Why does $\mathbf{U}^T\mathbf{U}$ cancel out)?
- Not difficult to see \mathbf{V} constitutes the principal components of \mathbf{X} .

Summary

- Utility matrices are a special kind of data matrices
 - Typically, you would want to change the representation of the data points among the feature dimension, but here both dimensions are important.
- We have seen two linear methods of dimensionality reduction commonly used in the recommender systems: PCA and SVD.
 - The crux of linear dimensionality reduction is finding new basis vectors.
 - PCA can be thought of changing the basis vector for the row/column space (depending on how you put the data) but SVD does that simultaneously.
- But they are on fully specified matrices: that's not the case for recommender systems:
 - Subtract row mean to convert a sparse utility matrix to a fully specified data matrix.

What is The Problem In Such Conversions?

- We have a data matrix X ($m \times n$), we want to reduce it to $m \times k$.
- If m is users and n is movies, it will give us a *compact* representation of users in movie space.
- We can use PCA.
- Also, to get the *compact* movie representation in terms of users, we can just transpose the matrix and run PCA again.
- SVD will do this jointly.

<i>User Index</i>	<i>Godfather</i>	<i>Gladiator</i>	<i>Nero</i>
1	1	1	1
2	7	7	7
3	3	1	1
4	5	7	7
5	3	1	?
6	5	7	?
7	3	1	?
8	5	7	?
9	3	1	?
10	5	7	?
11	3	1	?
12	5	7	?

- Correlation between Gladiator is Nero is extremely high.
- Correlation between Gladiator is Godfather is low.
- It's not possible to do PCA on this matrix because it is not fully specified.
- Let's make it fully specified by using the **average score** (column mean) for Nero which is 4.
- What happens to the covariance matrix?

	<i>Godfather</i>	<i>Gladiator</i>	<i>Nero</i>
<i>Godfather</i>	2.55	4.36	2.18
<i>Gladiator</i>	4.36	9.82	3.27
<i>Nero</i>	2.18	3.27	3.27

- Gladiator and Nero has low correlation, and Gladiator and Godfather has high correlation!

Latent Factor Models to the Rescue

- PCA could be thought as a matrix factorization:
- In general, we can factorize a data matrix.
 - You can factorize a matrix \mathbf{X} (m x n) in \mathbf{UV}^T .
 - \mathbf{U} has dimensions mxk and \mathbf{V} has dimensions nxk (if k is the rank of the matrix, such factorization will always exist).
 - U is a lower dimensional representation of X.
 - PCA can be thought of a factorization with some constraints: in U, covariances are minimized and variances are maximized.
 - SVD was $\mathbf{X} = \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T$. You can see by putting the values $\mathbf{\Sigma}$ inside \mathbf{U} or \mathbf{V} , this is again the same formulation. The constraints are \mathbf{U} and \mathbf{V} are orthonormal. Then we will take those \mathbf{U} and \mathbf{V} columns that corresponds to the highest k values of $\mathbf{\Sigma}$ (to get a k rank approximation of the original matrix)
- We will assume the factorization will discover the latent properties in the data (for movies-users, this can be the genre).
- How to find U and V? How to handle missing values?

PCA: we will find out a projection matrix \mathbf{P} such that $\mathbf{XP} = \mathbf{Y}$.

Original formulation was $\mathbf{X} = \mathbf{UV}^T$.

Not hard to see that $(\mathbf{V}^T)^{-1}$ would be \mathbf{P} .

Matrix Factorization

- A matrix \mathbf{X} ($m \times n$) with rank k ($k < \min(m, n)$) can always be factorized as $\mathbf{X} = \mathbf{UV}^T$ where \mathbf{U} is $m \times k$ and \mathbf{V} is $n \times k$.
- If we want a lower rank approximation ($< k$), we can still do the same, now $\mathbf{X} \approx \mathbf{UV}^T$
- How to find \mathbf{U} and \mathbf{V} ?
- How to handle missing values?
- Let's try an iterative process:
 - Start with random matrices \mathbf{U} and \mathbf{V} (\mathbf{UV}^T will obviously not equal \mathbf{X}).
 - At each step, calculate the approximation error: $J = 0.5 \|\mathbf{X} - \mathbf{UV}^T\|^2$ Frobenius norm
 - Depending on the error: *adjust* \mathbf{U} and \mathbf{V} .
 - Continue until E is less than a certain value.

What is the Error?

- $E = 0.5 ||\mathbf{X} - \mathbf{UV}^T||^2$
- Let's call \mathbf{UV}^T as \mathbf{D} . What is the dimension of \mathbf{D} ? $m \times n$.
- What is an element of \mathbf{D} ? $r'_{ij} = \sum_{s=1}^k u_{is} \cdot v_{js}$
- What is J then? $\frac{1}{2} \sum_{i=1}^m \sum_{j=1}^n (r_{ij} - \sum_{s=1}^k u_{is} v_{js})^2$
- But all values are not observed! in other words, you can not take that external sums.
 - Let's assume the set of observed values in \mathbf{X} is S .
 - Then the sum becomes: $\frac{1}{2} \sum_{(i,j) \in S} \left(r_{ij} - \sum_{s=1}^k u_{is} \cdot v_{js} \right)^2$
- If we can find u_{is} and v_{js} from this, we have solved the unseen ratings problem!

How to Adjust U and V?

- Gradient descent: $\theta_{new} = \theta_{old} + \alpha \nabla E$
- Suppose I have a parameter θ . (with a value)
- I will find the gradient of the error with respect to θ .
- My new value for the parameter will be in the direction of the gradient, multiplied with some constant.
- This will give me the new value for the parameters.
- What are my parameters here? U and V themselves, more specifically, the values u_{is} and v_{js} .
- remember, $J = \frac{1}{2} \sum_{(i,j) \in S} e_{ij}^2 = \frac{1}{2} \sum_{(i,j) \in S} \left(r_{ij} - \sum_{s=1}^k u_{is} \cdot v_{js} \right)^2$
- We will find out $\frac{\partial J}{\partial u_{iq}}$ and $\frac{\partial J}{\partial v_{jq}}$

Formal Algorithm for Matrix Factorization

- Let's look at the informal version first:

Start with random matrices \mathbf{U} and \mathbf{V} (\mathbf{UV}^T will obviously not equal \mathbf{X}).

At each step, calculate the approximation error: $J =$

$$0.5\|\mathbf{X}-\mathbf{UV}^T\|^2$$

Depending on the error: *adjust* \mathbf{U} and \mathbf{V} .

Continue until E is less than a certain value.

$$J = \frac{1}{2} \sum_{(i,j) \in S} e_{ij}^2 = \frac{1}{2} \sum_{(i,j) \in S} \left(r_{ij} - \sum_{s=1}^k u_{is} \cdot v_{js} \right)^2$$

$$\begin{aligned} \frac{\partial J}{\partial u_{iq}} &= \sum_{j:(i,j) \in S} \left(r_{ij} - \sum_{s=1}^k u_{is} \cdot v_{js} \right) (-v_{jq}) \quad \forall i \in \{1 \dots m\}, q \in \{1 \dots k\} \\ &= \sum_{j:(i,j) \in S} (e_{ij})(-v_{jq}) \quad \forall i \in \{1 \dots m\}, q \in \{1 \dots k\} \end{aligned}$$

$$\begin{aligned} \frac{\partial J}{\partial v_{jq}} &= \sum_{i:(i,j) \in S} \left(r_{ij} - \sum_{s=1}^k u_{is} \cdot v_{js} \right) (-u_{iq}) \quad \forall j \in \{1 \dots n\}, q \in \{1 \dots k\} \\ &= \sum_{i:(i,j) \in S} (e_{ij})(-u_{iq}) \quad \forall j \in \{1 \dots n\}, q \in \{1 \dots k\} \end{aligned}$$

Algorithm *GD*(Ratings Matrix: R , Learning Rate: α)

begin

Randomly initialize matrices U and V ;

$S = \{(i, j) : r_{ij} \text{ is observed}\}$;

while not(convergence) **do**

begin

Compute each error $e_{ij} \in S$ as the observed entries of $R - UV^T$;

for each user-component pair (i, q) **do** $u_{iq}^+ \leftarrow u_{iq} + \alpha \cdot \sum_{j:(i,j) \in S} e_{ij} \cdot v_{jq}$;

for each item-component pair (j, q) **do** $v_{jq}^+ \leftarrow v_{jq} + \alpha \cdot \sum_{i:(i,j) \in S} e_{ij} \cdot u_{iq}$;

for each user-component pair (i, q) **do** $u_{iq} \leftarrow u_{iq}^+$;

for each item-component pair (j, q) **do** $v_{jq} \leftarrow v_{jq}^+$;

Check convergence condition;

end

end

What are the Possible Improvements?

• Regularization

- Raters might have bias: some raters will always provide less score than the others.
- Movies might have bias: some movies will be popular across the board.
- These biases are unknown: you need to discover them from the data!
- Can we change the error function itself to capture this?

- Previously, $r'_{ij} = \sum_{s=1}^k u_{is} \cdot v_{js}$

- Now, $r'_{ij} = o_i + p_j + \sum_{s=1}^k u_{is} \cdot v_{js}$ o_i is the rater bias, p_j is the popularity bias.

- With that, we want to minimize: $\frac{1}{2} \sum_{i,j \in S} \left(r_{ij} - o_i - p_j - \sum_{s=1}^k u_{is} v_{js} \right)^2$

Note, o_i and p_j won't cancel out!

• Constraints

- If you cast SVD as a factorization problem, then you would want to have the constraints that \mathbf{U} and \mathbf{V} are orthonormal sets.

Summary

- Latent factor models work on a basic assumption: if you factorize a matrix and provide a low rank approximation, these low dimensional vectors would capture the latent reason for the data generation.

NERO	Julius Caesar	Patriot	Sleepless Seattle	Pretty Woman	Casablanca
1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	0	0	0
1	1	1	1	1	1
-1	-1	-1	1	1	1
-1	-1	-1	1	1	1
-1	-1	-1	1	1	1

User to genre mapping

history	romance
1	0
1	0
1	0
1	1
-1	1
-1	1
-1	1

Movie to genre mapping

history	1	1	1	0	0	0
romance	0	0	0	1	1	1

Discovery of latent factors is equivalent of finding new basis vectors.

This is what brings everything together,.

Summary

- A basic latent factor model is equivalent of linear dimensionality reduction techniques with some caveats.
 - Does not hold the provable guarantees of PCA.
 - The basis vectors are not orthonormal, hard to come to a nice solution mathematically (remember for PCA we had to compute the eigen vectors of a matrix?).
 - Therefore, needs to be computed through iterative method.
- However, they are more general.
 - With addition of constraints, most linear dimensionality reduction techniques can be cast as a latent factor model.
 - Also, solves the problem of missing values nicely!

DNNs for Recommender Systems: A Very Basic Intro

- What is a DNN good for?
 - Finding representation in an unsupervised way.
 - Fusing different modalities (text, image, ratings).
- In content-based recommender systems, you can use a DNN to find a k dimensional representation.
 - This is same as finding the representation for a sentence from text.
- You can also fuse multiple modalities:
 - Because DNNs represent each image/text as a vector, it's not difficult to combine them (typically with a concatenation but other functions can be used).
- Collaborative filtering can also be improved using DNNs.

Lecture/Course Summary

- This lecture
 - There is a unified view of matrix factorization, linear dimensionality reduction.
 - SVD and PCA are powerful tools, but understanding the underlying mathematics is also important.
 - This gives us the intuition to create new methods.
- Course:
 - **Sentiment Analysis:** What is the most complex part of the problem? Handling the context. "You are so bad, man!" . Also, going from word to sentence representation using ConvNet.
 - **Dimensionality reduction and representation:** Representation is challenging; nonlinear methods (Word2Vec) are often better than the linear ones but hard to interpret and comes with no guarantee. Linear methods (PCA) has a nice guarantee of residual error minimization and variance maximization: but does not perform so well on words.
 - **Recommender Systems:** Why they are important, SVD (joint optimization of basis vectors for row and column space), gradient descent (sort of).